

Hortonworks Tutorial 5

Optimization on Hive

In this tutorial, we'll focus on taking advantage of performance improvements of Hive on **TEZ**. We are going to look at some of the new features that HDP 2.5 brings to Apache Hive.

Task 1. Using HDP to upload files

We will use the same dataset used in Tutorial 1. If you do not have the files, download the files from this link: <http://opensourceports.com/files/basketball/BasketballDB-20130121.zip>. We will use the following 3 files: `basketball_master.csv`, `basketball_players.csv` and `basketball_teams.csv`. Make sure these files are uploaded to HDFS.

Select the "Hive View" menu. In Tutorial 1, you created three tables of `basketball_teams`, `basketball_players`, and `basketball_master`. You will use those tables in this tutorial. This means you should see these 3 tables in **Database Explorer**.

You can view the 100 sample records when clicking on the "Load sample data" (e.g. the one highlighted in yellow) icon right besides each table as seen in Fig. 1:

The screenshot displays the Hortonworks Data Platform interface. On the left, the 'Database Explorer' shows a list of databases, with 'basketball_master' highlighted. The main area shows a SQL query editor with the query: `1 SELECT * FROM basketball_master LIMIT 100;`. Below the query editor, the 'Query Process Results' section shows the status 'SUCCEEDED' and a table of sample data.

basketball_master.bioID	basketball_master.firstName	basketball_master.lastName
bioID	firstName	lastName
bioID	firstName	lastName
abdela01	Alaa	Abdelnaby
abdulka01	Kareem	Abdul-Jabbar
ahdulma01	Mahdi	AhduL-Dahman

Figure 1: Sample data in the table, `basketball_master`

Task 2. Speed Improvements

1. To take a look at the speed improvements of Hive on **TEZ**, we can run some sample queries. For this we will use the above three tables.
2. By default, the Hive view runs with **TEZ** as its execution engine. That is because **TEZ** has great speed improvements over the original MapReduce execution engine. But how much

are exactly these improvements? Well, let's find out! Click on the "Settings" tab (see "1" in Fig. 2). Then click on "+Add" (see "2" in Fig. 2).

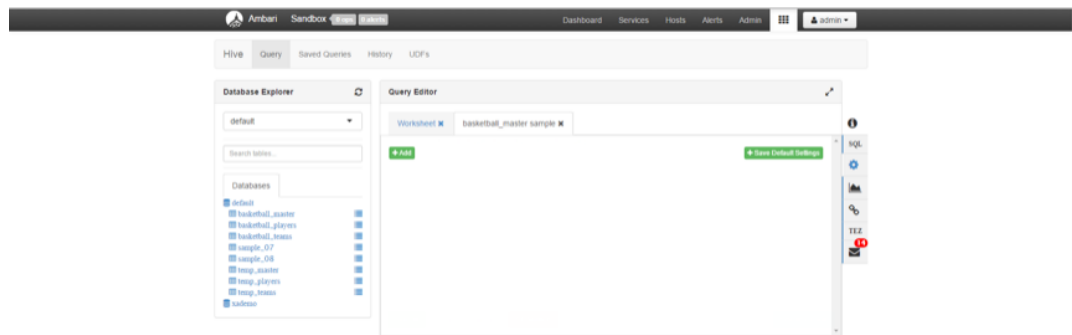


Figure 2: The "Settings" menu in Hive View

3. Try to find the property which is "hive.execution.engine". Select this property and then for its value select "mr" (short for MapReduce) (see Fig. 3).

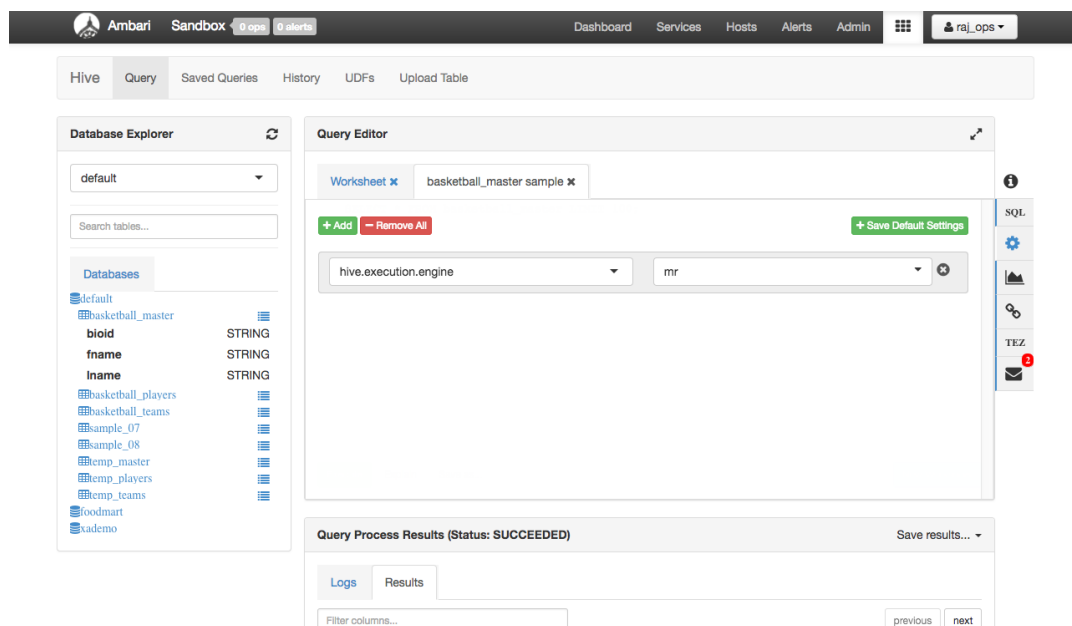


Figure 3: Using MapReduce framework

4. We now test a query using the MapReduce framework. Execute the following query and wait for the results:

```
select m.*, p.playerID, p.year, p.points from basketball_master m
join basketball_players p on m.bioID = p.playerID;
```

This query is run using the MapReduce framework as seen in "Logs" in Fig. 4 (see the last 5 lines):

The screenshot displays the Hive web interface. On the left, a sidebar shows a list of databases including 'default', 'basketball_master', 'basketball_players', 'basketball_teams', 'sample_07', 'sample_08', 'temp_master', 'temp_players', 'temp_teams', 'foodmart', and 'demo'. The main area shows a worksheet titled 'basketball_master sample' with the following SQL query:

```
1 select m.*, p.playerID, p.year, p.points
2 from basketball_master m join basketball_players p
3 on m.bioid = p.playerID;
```

Below the query, there are buttons for 'Execute', 'Explain', 'Save as...', and 'New Worksheet'. The 'Execute' button has been clicked, and the 'Query Process Results (Status: SUCCEEDED)' are displayed. The results are shown in a tabbed interface with 'Logs' and 'Results' tabs. The 'Logs' tab is active, showing the following log output:

```
INFO : Execution completed successfully
INFO : MapredLocal task succeeded
INFO : Number of reduce tasks is set to 0 since there's no reduce operator
INFO : number of splits:1
INFO : Submitting tokens for job: job_1492572442634_0001
INFO : The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1492572442634_0001/
INFO : Starting Job = job_1492572442634_0001, Tracking URL =
http://sandbox.hortonworks.com:8088/proxy/application_1492572442634_0001/
INFO : Kill Command = /usr/hdp/2.5.0.0-1245/hadoop/bin/hadoop job -kill job_1492572442634_0001
INFO : Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
INFO : 2017-04-19 03:40:55,173 Stage-3 map = 0%, reduce = 0%
INFO : 2017-04-19 03:41:01,525 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 3.72 sec
INFO : MapReduce Total cumulative CPU time: 3 seconds 720 msec
INFO : Ended Job = job_1492572442634_0001
```

Figure 4: Logs using the MapReduce framework

5. Now instead of using the MapReduce framework we enable TEZ. Click on the Settings in the Hive view, and then select "hive.execution.engine" and "tez". Execute the same query as we had run earlier to see the speed improvements with TEZ. The printed log should be different to the log created when using MapReduce, and the time taken should be reduced. See Fig. 5.

The screenshot displays the Databricks SQL interface. On the left, a sidebar shows a list of databases including 'default', 'basketball_master', 'basketball_players', 'basketball_teams', 'sample_07', 'sample_08', 'temp_master', 'temp_players', 'temp_teams', 'foodmart', and 'xademo'. The main workspace area shows a worksheet titled 'basketball_master sample' with a SQL query:

```
1 select m.*, p.playerID, p.year, p.points
2 from basketball_master m join basketball_players p
3 on m.bioID = p.playerID;
```

Below the query editor, there are buttons for 'Execute', 'Explain', 'Save as...', and 'New Worksheet'. The 'Query Process Results (Status: SUCCEEDED)' section is expanded, showing the 'Logs' tab. The log content is as follows:

```
INFO : Tez session hasn't been created yet. Opening session
INFO : Dag name: select m.*, p.playerID, p.year,...p.playerID(Stage-1)
INFO : Setting tez.task.scale.memory.reserve-fraction to 0.30000001192092896
INFO :
INFO : Status: Running (Executing on YARN cluster with App id application_1492572442634_0002)

INFO : Map 1: -/- Map 2: -/-
INFO : Map 1: 0/1 Map 2: 0/1
INFO : Map 1: 0(+1)/1 Map 2: 0(+1)/1
INFO : Map 1: 0(+1)/1 Map 2: 0(+1)/1
INFO : Map 1: 1/1 Map 2: 0(+1)/1
INFO : Map 1: 1/1 Map 2: 1/1
```

Figure 5: Logs using TEZ

6. Now let's try the following new query to work with. First execute the query first using "mr", then using "tez". The completion time should be significantly reduced when "tez" is used. Of course, the difference might be differently measured depending on the performance of the machine you use.

```
select m.fname, m.lname, sum(p.points) as totalPoints
from basketball_master m
join basketball_players p on m.bioID = p.playerID
group by m.fname, m.lname order by totalPoints desc;
```

7. When TEZ is enabled, you can also click "TEZ" tab on the right to see the DAG detail including the time taken, which is not available if executing Hive on "mr". See Fig. 6.
8. You can even obtain the time taken for each Map job by clicking on the "All Vertices" or "All Tasks" tab (e.g. Fig. 7 shows what happens when clicking on "All Vertices").

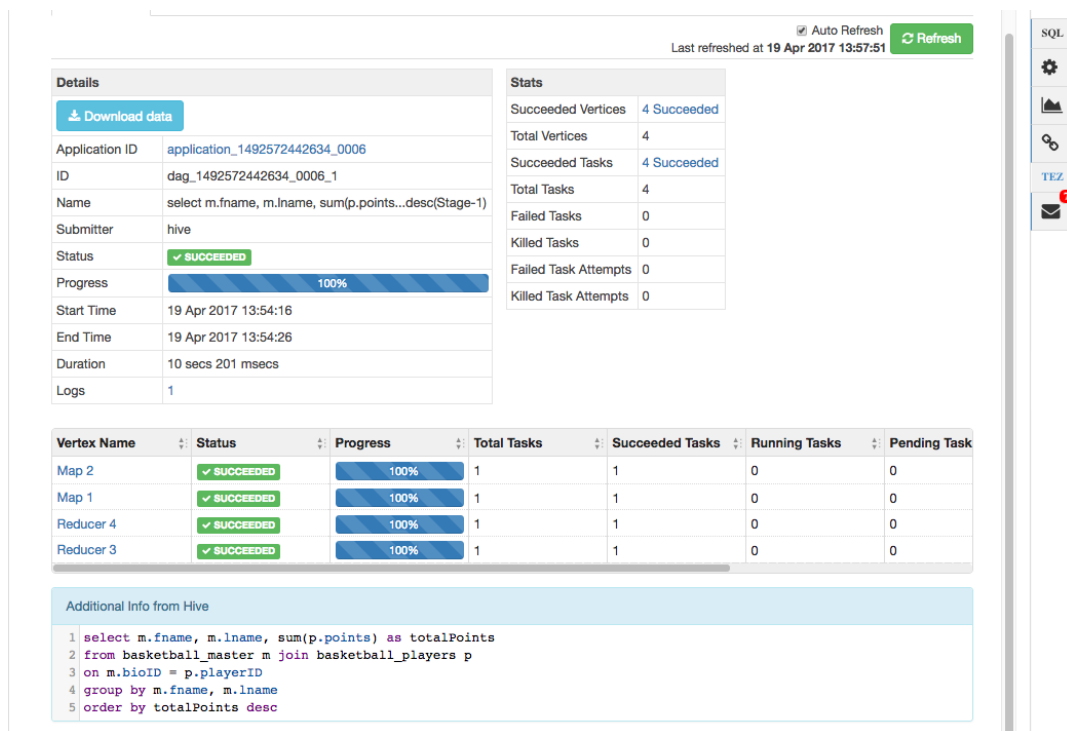


Figure 6: DAG Details in TEZ

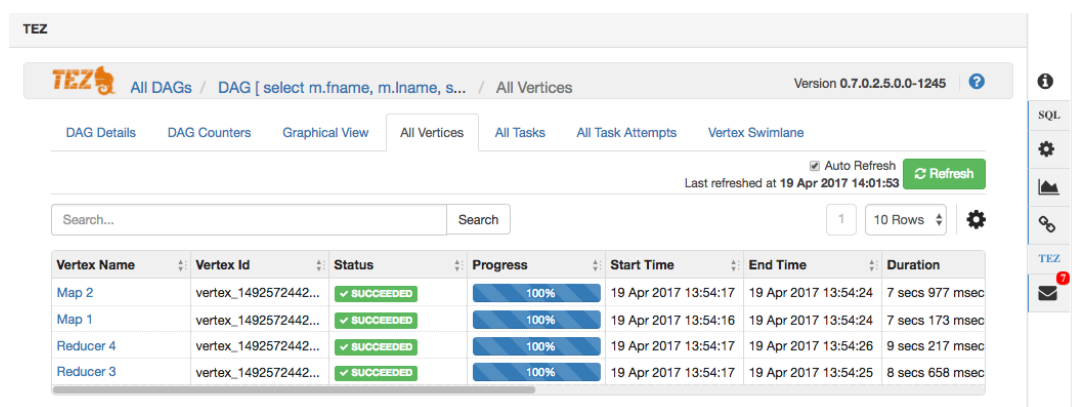


Figure 7: The "All Vertices" tab

- Try clicking on the different parts above, such as "Graphical View" and explore some of the other execution information from Tez (see Fig. 8).

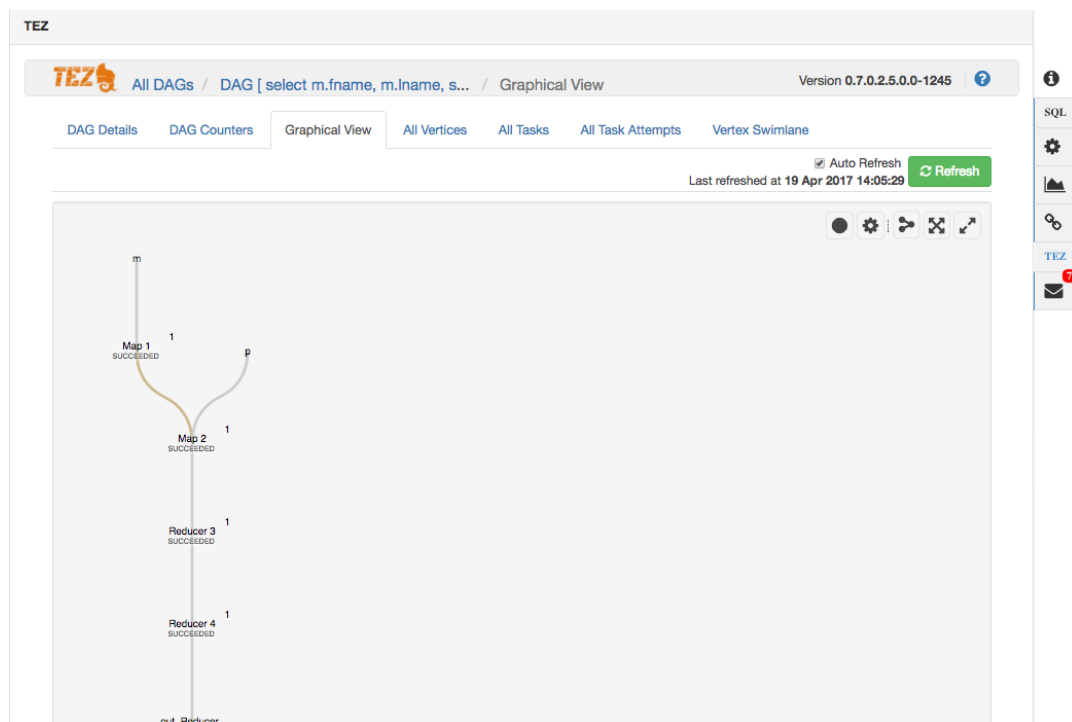


Figure 8: The "Graphical View" tab

10. You can obtain further details about Tez and MapReduce jobs in the HDP Sandbox Web UI. If you go to this page <http://127.0.0.1:8088/cluster>, you can track your jobs. You can click on your job and see further details.

ID	User	Name	Application Type	Queue	Application Priority	Start Time	Finish Time	State	Final Status	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	% of Queue	% of Cluster	Progress
application_1492572442634_0006	hive	HIVE-40fc0e90-293e-41ba-af20-d81077837ba9	TEZ	default	0	Wed Apr 19 13:54:10 +1000 2017	Wed Apr 19 14:05:20 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	
application_1492572442634_0005	hive	select m.fname, m.lname, sum(p.points_desc(Stage-3))	MAPREDUCE	default	0	Wed Apr 19 13:53:05 +1000 2017	Wed Apr 19 13:53:22 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	
application_1492572442634_0004	hive	select m.fname, m.lname, sum(p.points_desc(Stage-2))	MAPREDUCE	default	0	Wed Apr 19 13:52:46 +1000 2017	Wed Apr 19 13:53:03 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	
application_1492572442634_0003	hive	HIVE-86d4f4f8-66bf-41d9-b098-0f27ee55fb5a	TEZ	default	0	Wed Apr 19 13:51:53 +1000 2017	Wed Apr 19 14:03:03 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	
application_1492572442634_0002	hive	HIVE-c42d78c-3713-4dc3-3b81-0945d5d447aa	TEZ	default	0	Wed Apr 19 13:46:12 +1000 2017	Wed Apr 19 13:57:22 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	
application_1492572442634_0001	hive	select m.*, p.playerID, ...	MAPREDUCE	default	0	Wed Apr 19 13:46:12 +1000 2017	Wed Apr 19 13:57:22 +1000 2017	FINISHED	SUCCEEDED	N/A	N/A	N/A	0.0	0.0	

Figure 9: Tracking jobs

Task 3. Cost Based Optimization (CBO)

Cost Based Optimisation (CBO) engine uses statistics to optimize query processing in Hive. In the new version of Hortonworks, CBO is enabled by default. You can check this by selecting Ambari → Hive → Configs (as shown in Fig. 10).

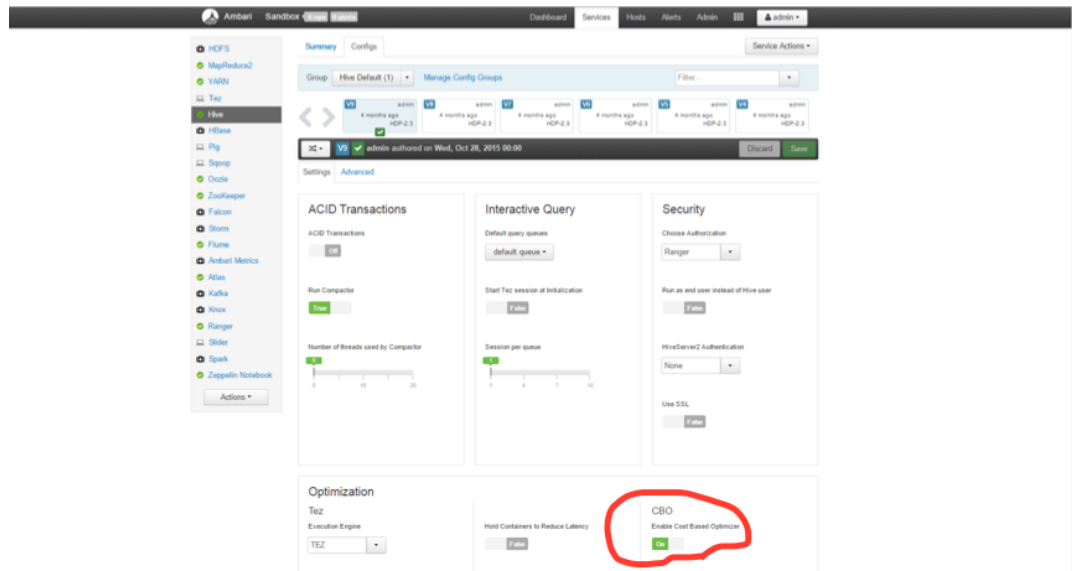


Figure 10: CBO Config

While CBO is enabled in Hive, it is not utilised for optimization due to missing statistics about tables and columns but by providing these statistics, it can be used, specially when there are joins between multiple tables.

1. **Test with queries:** Let's do a simple exercise. Let's run the following query in which on theSettingsin the Hive view, you need to choose "hive.execution.engine" and "tez".

```
SELECT a.year, a.playerID, c.fname, c.lname, a.points
FROM basketball_players a JOIN
(SELECT year, max(points) points FROM basketball_players GROUP BY year) b
ON (a.year = b.year AND a.points = b.points) JOIN basketball_master c
ON (c.bioID = a.playerID) ORDER BY year DESC;
```

Now, click on "Explain" for the above query. The result in the figure below shows the CBO was not used because of missing statistics (see Fig. 11).

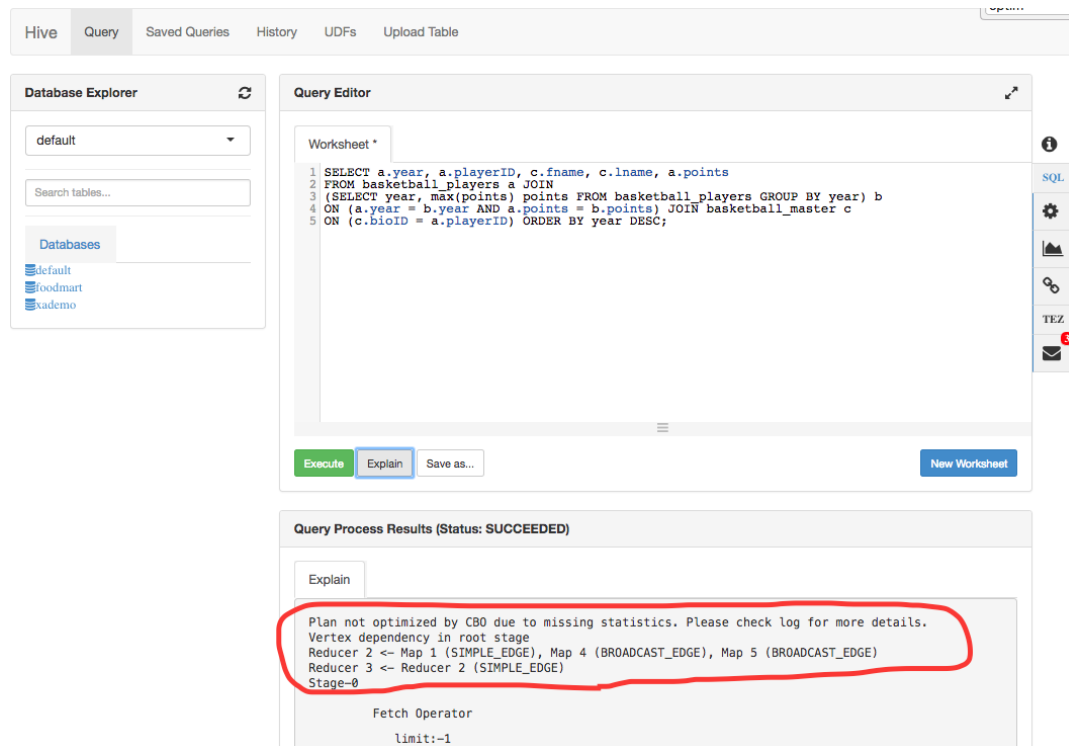


Figure 11: Not optimized result by CBO

To provide the missing statistics, write and execute the following commands ONE by ONE (it will give an error if you execute them together).

- *ANALYZE TABLE basketball_players COMPUTE STATISTICS;*
- *ANALYZE TABLE basketball_players COMPUTE STATISTICS FOR COLUMNS year, playerID, points;*
- *ANALYZE TABLE basketball_master COMPUTE STATISTICS;*
- *ANALYZE TABLE basketball_master COMPUTE STATISTICS FOR COLUMNS bioID, fname, lname;*

Now you can execute the query we tested earlier and after getting the results, click on Explain (see Fig. 12).

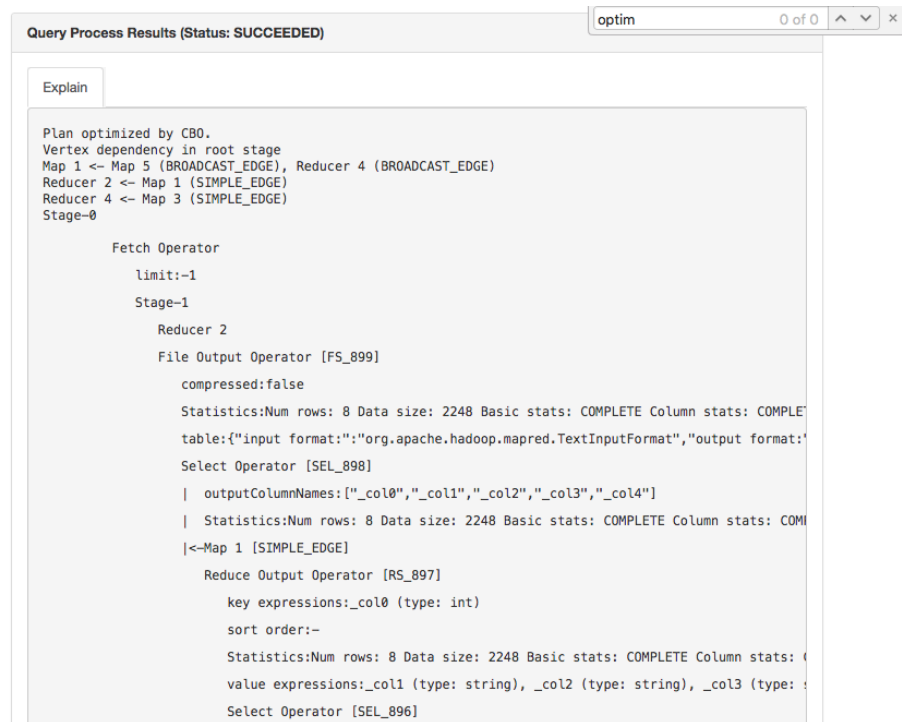


Figure 12: Optimized result by CBO

2. Visit and investigate following links to more fully understand TEZ and CBO:

- <http://hortonworks.com/hadoop-tutorial/supercharging-interactive-queries-hive-tez/>
- <http://hortonworks.com/blog/hive-0-14-cost-based-optimizer-cbo-technical-overview/>